



# Parallelized Neural Network training with the MNIST dataset

The RISCtakers

---

# Background



# What is a neural network?

- A type of an AI model inspired by the workings of the human brain
- Designed to recognize patterns in order to problem solve
- Consists of a network of interconnected nodes that process data
- Consists of an input layer, one or more hidden layers, and an output layer
- Learns through training, getting smarter with the goal of minimizing the error
  - Through forward propagation, backpropagation, and gradient descent and loss functions

# MNIST dataset

- Contains 70,000 images of handwritten digits (0 through 9)
- Each image is a bitmap consisting of 28x28 pixels, representing a single digit
- Training set is 60,000 images, with a test set of 10,000 images
- The problem's simplicity creates a good environment to evaluate machine learning techniques and algorithms




---

# How does a neural network work?

# Layers

- A neural network is composed of three layers





# Input layer

- The starting point of the neural network
- Receives the input data, each neuron represents one feature of the input data
- Each neuron connects to every other neuron
- Each neuron represents one feature in the image
- In the context of the MNIST data set, each neuron represents the intensity of each grayscale pixel in the image.



# Hidden layers

- Where the majority of the computation is done
- “Hidden” means that they aren’t directly in contact with the input/output of the neural network
- Each neuron performs a weighted sum of inputs from the previous layer, applies a bias, then sends to an activation function
  - Non-linearity helps the network learn more complex relationships
- Different hidden layers can have different functions
  - Early layers might detect simple color features and shapes
  - Deeper layers may identify more complex patterns or shapes
- The number of hidden layers depends on how complex the task is
  - MNIST can be processed with just 3, may be improved with N layers.





# Output layer

- Final layer in the network
  - Presents the results of all of the network's computations
- These neurons represent the final output of the network
  - Output depends on the task
  - Could be a classification table, a single value, or in our case a 10 element vector with the probability that the input image is a certain digit
- [0.1, 0.05, **0.6**, 0.1, 0.02, 0.04, 0.01, 0.04, 0.1, 0.2]
  - 60% chance that the image is the digit '2'
- The output is a probability distribution showing the network's confidence level for each image



# Layer Initialization

```
NeuralNetwork::NeuralNetwork(int inputSize, int hiddenSize, int outputSize, double learningRate, int numLayers)
{
    // add input layer to network
    network_layers.emplace_back(Layer(inputSize, hiddenSize));

    // add hidden layers
    for (int i = 0; i < numLayers; i++) {
        network_layers.emplace_back(hiddenSize, hiddenSize);
    }

    // add output layer
    network_layers.emplace_back(hiddenSize, outputSize);
}
```

---


# Approach



## Optimization method: Stochastic gradient descent

- This is an optimization algorithm used to minimize the loss function, updating the model's weights and biases accordingly
- Uses small batch of samples to compute the gradient
- In the context of MNIST, SGD works by dividing the process into small batches
  - It makes predictions based on the current weights and biases
  - The loss is calculated using these predictions and the correct labels of the batch
  - The gradient is computed with respect to the loss of each parameter
  - The weights and biases are updated in the opposite direction of the gradient
- More efficient than gradient descent, which looks at the dataset as a whole.

# SGD visualization





# Forward propagation

For each layer in the neural network, call the forward method and return the newly formulated vector.


Foreshadowing! ->

```
vector<double> NeuralNetwork::forward(vector<double> input) {  
    // copy data  
    vector<double> output = input;  
  
    for (auto &layer: network_layers) {  
        output = layer.forward_propagation_serial(output);  
        // output = layer.forward_propagation_parallelized(output);  
        // output = layer.forward_propagation_mkl(output);  
    }  
  
    return output;  
}
```

# Backpropagation

- The process of computing gradients of the loss in respect to the weights and biases
- Propagates this gradient backward through the layers of the network, layer by layer
- Training stops when a certain number of epochs is reached, or improvement is no longer substantial
- Key component to neural network training

Training in progress. . .



*Credit: 3Blue1Brown*



# Loss function

Crucial for gradient descent optimization. Compares the the actual to the target to determine how much a feature must be changed.

```
double NeuralNetwork::computeLoss(const vector<double> output, const vector<double> target)
{
    double loss = 0.0;

    // square loss function
    for (int i = 0; i < outputSize; ++i)
    {
        loss += pow(target.at(i) - output.at(i), 2);
    }
    return loss / output.size(); // Returns the average loss
}
```



---

**Why can this be parallelized?**



# What makes this a parallelization problem?

1. Forward and backward propagations rely on vector multiplication, with layer of neurons being updated in an embarrassingly parallel manner.
2. Weights and biases updates
  - a. After these are computed, they do need to be synchronized, but can often be done in parallel for different parts of the network
3. MNIST is a reasonably large dataset, and processing the data in parallel can lend a significantly faster training time



## What makes this a parallelization problem? (2)

- Vector Multiplication
  - Vector multiplication in forward propagation
  - Each element computed independently, inherently parallelizable
  - Vectors are multiplied by the weight matrix to produce another vector, then passed into an activation function. This process continues through multiple layers.
- Hardware optimization
  - Running this neural network model on a GPU would take advantage of its parallel nature. Optimized for fast vector and matrix operations
  - Hardware acceleration is one of the reasons why neural networks have become feasible for deep learning models.

---

# Vector Multiply Approaches



## Serial Multiply

```
std::vector<double> Layer::forward_propagation_serial(const std::vector<double> &input_vector) {  
    std::vector<double> output_vector;  
    double value;  
  
    for (int i = 0; i < outputs; i++) {  
        value = 0;  
        for (int j = 0; j < input_vector.size(); j++) {  
            value += input_vector.at(j) * weight_matrix.at(i).at(j);  
        }  
        output_vector.at(i) = sigmoid(value + bias);  
    }  
    return output_vector;  
}
```



# Parallel For

```
std::vector<double> Layer::forward_propagation_parallelized(const std::vector<double> &input_vector) {  
    std::vector<double> output_vector;  
    double value;  
  
    #pragma omp parallel for  
    for (int i = 0; i < outputs; i++) {  
        value = 0;  
        for (int j = 0; j < inputs; j++) {  
            value += input_vector.at(j) * weight_matrix.at(i).at(j);  
        }  
        output_vector.at(i) = sigmoid(value + bias);  
    }  
    return output_vector;  
}
```



# Intel MKL

```
std::vector<double> Layer::forward_propagation_mkl(const std::vector<double> &input_vector) {  
    std::vector<double> output_vector;  
  
    // do vector multiplication  
    cblas_dgemv(CblasRowMajor, CblasNoTrans,  
        inputs, outputs, 1.0, weight_matrix.data(), inputs, input_vector.data(), 1, 1.0, output_vector.data(), 1);  
    return output_vector;  
}
```



# Challenges to look out for

- Hyperparameter tuning
  - Being sure to choose the correct learning rate, number of layers, number of neurons in each layer
  - Inappropriate choices may lead to a poor learning rate or slow convergence
- Model explainability
  - Due to the complex nature of a neural network, it can be challenging to understand why a certain prediction is happening
  - Many parameters and different possible outputs
  - Ensuring that a model is consistent across tests and the data can be repeated is important for creating a strong neural network



---

**Thanks!**